

---

# REPOEXEC: Evaluate Code Generation with a Repository-Level Executable Benchmark

---

Nam Le Hai\*, Dung Manh Nguyen\*, Nghi D. Q. Bui\*,†

\* FPT Software AI Center, Viet Nam

†Fulbright University, Viet Nam

namlh35@fpt.com, dungnm31@fpt.com, bdqnghi@gmail.com




Figure 1: Data Processing Pipeline of REPOEXEC

## Abstract

The ability of CodeLLMs to generate executable and functionally correct code at the *repository-level scale* remains largely unexplored. We introduce REPOEXEC, a novel benchmark for evaluating code generation at the repository-level scale. REPOEXEC focuses on three main aspects: executability, functional correctness through automated test case generation with high coverage rate, and carefully crafted cross-file contexts to accurately generate code. Our work explores a controlled scenario where developers specify necessary code dependencies, challenging the model to integrate these accurately. Experiments show that while pretrained LLMs outperform instruction-tuned models in correctness, the latter excel in utilizing provided dependencies and demonstrating debugging capabilities. We also introduce a new instruction-tuned dataset that focuses on code dependencies and demonstrate that CodeLLMs fine-tuned on our dataset have a better capability to leverage these dependencies effectively. REPOEXEC aims to provide a comprehensive evaluation of code functionality and alignment with developer intent, paving the way for more reliable and applicable CodeLLMs in real-world scenarios. The dataset and source code can be found at <https://github.com/FSoft-AI4Code/RepoExec>.

## 1 Introduction

Code Large Language Models (CodeLLMs) have emerged as a prominent solution to assist with coding tasks recently Wang et al. [2021, 2023], Feng et al. [2020], Allal et al. [2023], Li et al. [2023], Lozhkov et al. [2024], Guo et al. [2024], Pinnaparaju et al. [2024], Zheng et al. [2024], Roziere et al. [2023], Nijkamp et al. [2022], Luo et al. [2023], Xu et al. [2022], Bui et al. [2023]. While CodeLLMs are adept at capturing the immediate context of the code being written, they often do not fully exploit the broader context available within the entire code repository. This limitation stems from the token length restrictions of general LLMs, which also apply to CodeLLMs. Consequently, suggestions generated by these models may not be optimally aligned with the project’s architecture or intended

functionality, as they tend to overlook the rich information embedded in related files, class hierarchies, dependencies, and more. Therefore, evaluating the ability of CodeLLMs on code generation at the repository-level scale is crucial to accurately measure their performance in real-world scenarios. There are existing repository-level code generation benchmarks, such as RepoBench Liu et al. [2023b], RepoCoder Zhang et al. [2023b], CrossCodeEval Ding et al. [2023], CoCoMIC Ding et al. [2024], and DevEval Li et al. [2024]. Although these benchmarks provide a comprehensive perspective on evaluating code generation at the repository-level scale, they lack 3 major features that prevent them from accurately reflecting real-world coding scenarios. First, code generation in the real world requires the code to be executable, and none of the existing benchmarks focus on this aspect. Second, these benchmarks rely on simple methods to retrieve contexts, such as including all imports and function calls (RepoBench) or simply retrieving similar files as contexts (RepoCoder), which might introduce noise into the contexts and impact the performance of code generation systems on these benchmarks. Third, generating real world code at repository-level scale requires the CodeLLMs to be able to capture contexts (both cross-file and in-file effectively), which may pose several questions such as if the dependencies provided by human are really the dependencies that the LLM need, or how well the CodeLLMs process the dependencies given that lots of dependencies feed into the prompt, do they use all of them or skip some of them?

To effectively address these shortcomings, we introduce REPOEXEC, a novel benchmark for evaluating code generation at the repository level that emphasizes executability and correctness. This benchmark is specifically designed to overcome the limitations of existing benchmarks by focusing on real-world applicability and a comprehensive evaluation of code functionality with 3 main focus:

- **Enhanced Executability:** REPOEXEC ensures that generated code is not only syntactically correct but also fully executable in a real-world environment by establishing an automated system that verifies installation and runtime requirements.
- **Dynamic Test Case Generation:** REPOEXEC incorporates a sophisticated mechanism for dynamically generating high-coverage test cases tailored to the functionality of the newly generated code, ensuring that the code performs its intended tasks accurately.
- **Dependency Usage:** REPOEXEC provides a pipeline to evaluate the usage of dependencies across CodeLLMs, offering insights on how well code dependencies should be provided. Additionally, we introduce an instruction-tuning dataset with dependency calls, which significantly improves CodeLLMs’ ability to leverage code dependencies and produce better results.

In addition, we introduce the *Dependency Invocation Rate (DIR)*, a novel metric that measures the proportion of code dependencies successfully integrated into the generated code. Higher DIR values indicate better model comprehension of dependencies, while lower values suggest challenges in dependency utilization.

Our experiments with REPOEXEC reveal key insights into repository-level code generation by CodeLLMs. Models achieve optimal performance with full context dependencies, though smaller contexts outperform medium ones due to formatting issues in the BasePrompt. While pretrained LLMs like Codellama-34b-Python lead in pass@1 rates, instruction-tuned LLMs is better at managing dependencies, despite sometimes generating overly complex code. In contrast, pretrained LLMs produce accurate code but struggle with dependency utilization. Multi-round debugging tests, especially with GPT-3.5 and WizardCoder, significantly improve both pass@1 scores and DIR after three rounds. Fine-tuning with our dependency-enhanced dataset not only boosts pass@1 and DIR metrics but also minimizes computational costs, highlighting the importance of executable code testing and the advantages of instruction-tuning and multi-round debugging in enhancing CodeLLMs’ performance and dependency management. In summary, our contributions are as follows:

1. We introduce REPOEXEC, a novel benchmark designed to evaluate code generation at the repository level with a focus on executability and correctness. This benchmark addresses the gaps in existing systems by emphasizing real-world applicability and providing a comprehensive assessment of code functionality.
2. REPOEXEC ensures that generated code is not only syntactically correct but also fully executable in real-world environments. This is achieved through an automated system that verifies installation and runtime requirements, ensuring the code can function in actual software settings.

3. Our benchmark includes a sophisticated mechanism for dynamically generating high-coverage test cases that are tailored to the functionality of the newly generated code. This ensures the code not only runs but also accurately performs its intended tasks.
4. We implement a pipeline within REPOEXEC to assess how effectively CodeLLMs utilize code dependencies. This includes an novel instruction-tuning dataset that enriches CodeLLMs’ capability to leverage these dependencies for improved code generation.
5. We introduce the Dependency Invocation Rate (DIR), a novel metric that measures the proportion of provided dependencies successfully incorporated into the generated code. This metric helps gauge the models’ understanding and utilization of dependencies.
6. Our experiments provide crucial insights into the performance of CodeLLMs in repository-level code generation. We demonstrate that while foundation models achieve high initial accuracy, instruction-tuned models excel in dependency management and addressing complex scenarios. Additionally, our multi-round debugging tests show notable improvements in model performance, particularly with enhancements in dependency management.

## 2 Related works

Coding-related tasks have been crucial for assessing the performance of Large Language Models (LLMs), with code generation emerging as a primary focus [Chen et al., 2021, Li et al., 2023, Jiang et al., 2024, Touvron et al., 2023, Roziere et al., 2023, Xu et al., 2022, Allal et al., 2023, Nijkamp et al., 2022]. Early benchmarks have been introduced to address this issue Yin et al. [2018], Iyer et al. [2018], Manh et al. [2023], Chen et al. [2021], Austin et al. [2021], Hendrycks et al. [2021]; however, they often had limited scope or employed weak evaluation approaches. Some benchmarks Yin et al. [2018], Iyer et al. [2018], Manh et al. [2023] exhibit domain diversity akin to real-world applications; however, their evaluation methodologies are constrained to match-based metrics, thereby decreasing the reliability of these benchmarks Chen et al. [2021]. Meanwhile, benchmarks with reliable evaluation approaches such as HumanEval Chen et al. [2021], MBPP Austin et al. [2021] and APPS Hendrycks et al. [2021] often entail limitations to specific domains like competitive programming. Besides, all the benchmarks mentioned primarily focus on standalone function generation, lacking consideration for cross-contextual and dependency invocation scenarios.

Several recent studies have introduced frameworks and benchmarks for repository-level code generation [Ding et al., 2024, Shrivastava et al., 2023, Ding et al., 2023, Liao et al., 2023, Liu et al., 2023b]. These studies closely align with real-world scenarios, underscoring the importance of cross-contextual considerations. However, these works are still limited to match-based evaluation methods. A recent study Li et al. [2024] introduced the DevEval benchmark for code generation within repository contexts, evaluating performance based on extracted tests available in the repository. These benchmarks primarily focus on assessing the functional correctness of generated outputs but have not extensively investigated the correctness in conjunction with the ability to utilize dependency contexts.

## 3 Data Collection Pipeline

### 3.1 Data Source

Developing an executable benchmark within repository contexts is challenging due to complex setup requirements and frequent lack of clear installation guidelines in repositories. Previous studies [Ding et al., 2024, Shrivastava et al., 2023, Ding et al., 2023, Liao et al., 2023] have often relied on match-based metrics for evaluation, which may not fully capture code functionality. In addition, test cases are essential for assessing code functionality. However, extracting test cases from repositories [Zhang et al., 2023a, Li et al., 2024, Zhang et al., 2024] often relies on available functions and heuristic rules, limiting adaptability and excluding data without existing tests. For instance, Li et al. 2024 found that over 99% of functions were discarded due to the absence of suitable tests. We propose a dataset collection pipeline to ensure that repositories can build executable environments and that test cases are automatically generated.

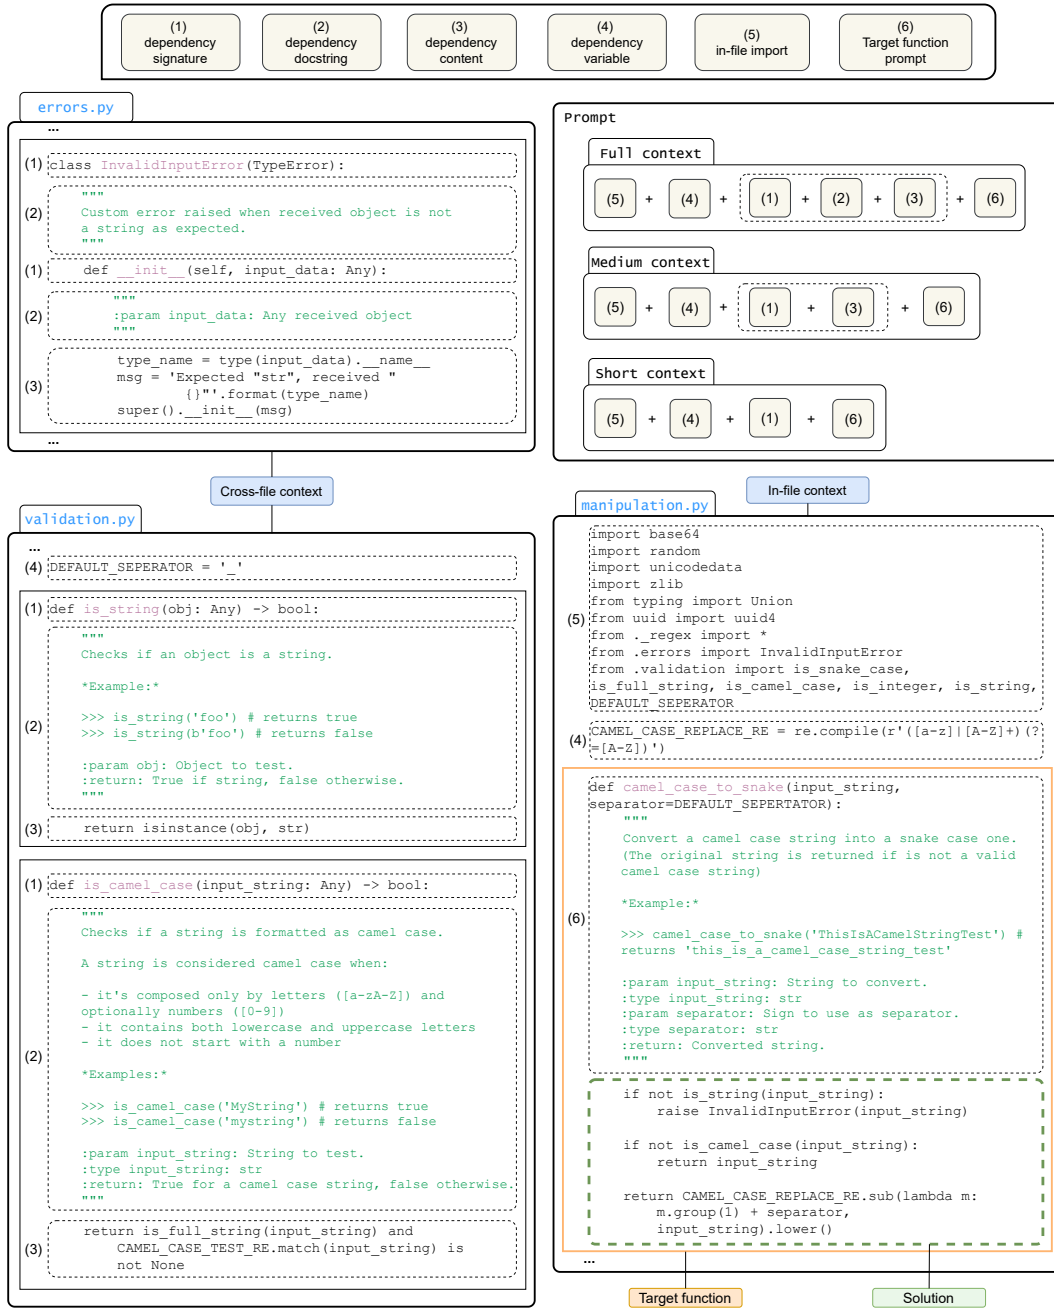


Figure 2: Illustration of a data instance in RepoExec. The target function signatures and their associated docstrings, which describe the functionality of the functions, are shown in (6). The in-file-imports and variable declarations are represented by (5) and (4), respectively. The remaining components, (1), (2), and (3), represent the function and class contexts. Specifically, (1) denotes the class or function signature, (2) may contain the description of the class, and (3) represents the function body of the cross-file function.

### 3.2 Functions and Dependencies Extraction

We extract functions and their dependencies from repositories, considering only those suitable for function-level code generation, akin to benchmarks like HumanEval Chen et al. [2021] and MBPP Austin et al. [2021]. Using tree-sitter, we parse files into Abstract Syntax Trees (AST) to extract functions, focusing on those with comprehensive docstrings and excluding functions used as entry points or that do not produce verifiable outputs.

**Dependencies extraction:** We employ static analysis to identify dependencies, excluding identifiers that are function parameters or typing objects. Each dependency name is then mapped to its definition within the repository using a repository graph and static analysis tools.

For example in Figure 2, for the function `camel_case_to_snake` in `manipulation.py`, our process identifies `CAMEL_CASE_REPLACE_RE` as an in-file dependency and analyzes import statements to track cross-file dependencies from `errors.py` and `validation.py`. Dependencies are then parsed and incorporated into the input prompt to ensure comprehensive context for code generation.

### 3.3 Test case generation

To overcome the limitation of requiring available tests in the repository for evaluation (as mentioned in Section 3.1), we leverage large language models (LLMs) to generate test cases automatically. Our proposed approach ensures the correctness of created test cases while also controlling and enhancing test coverage. To execute and validate the generated test cases, it is necessary to configure each repository to create an executable environment. Follow Lemieux et al. 2023, we use `pipreqs`<sup>1</sup> to identify each project’s dependencies.

In our test generation process, we conduct two phases corresponding to **Correctness Control** and **Coverage Enhancement**. After the test generation process, we exclude samples with a coverage percentage lower than 40%, as they are of insufficient quality to accurately assess the correctness of the generated code.

#### 3.3.1 Correctness Control

In this stage, we present the procedure for generating initial test cases and ensure the tests’ correctness. Specifically, we use CodeLlama-13b Roziere et al. [2023] and provide the model with the prompt detailed in Appendix C. After generation, we extract the first 20 assertion statements to construct the test cases. To ensure their correctness, we employ syntax and execution-based filtering methods to exclude low-quality test cases.

**Syntax-based filter:** We filter out tests that present syntax errors during parsing into AST. Additionally, tests that do not invoke the function under test are excluded. For instance, while the statement “`assert 1`” may pass during execution, it is meaningless and negatively impacts the evaluation.

**Execution-based filter:** We use `pytest`<sup>2</sup> to execute the generated tests. Tests are discarded if their execution output displays an error, except for `AssertionError`. If an assertion error is observed, we try to fix it using the Assertion Fixer.

- **Assertion Fixer:** The fixer resolves this error by executing the test case with the target function, extracting the output, and reassigning it to the assertion statement. Given that the return type can be complex in the repository (e.g., a defined object), we use `pickle`<sup>3</sup> to preserve the execution results.
- **Multiple time execution:** We witness some flaky tests, which produce inconsistent outcomes upon multiple executions due to the inherent randomness in their implementation. To eliminate these instances, we execute each test 10 times to compare the outputs.

---

<sup>1</sup><https://github.com/bndr/pipreqs>

<sup>2</sup><https://github.com/pytest-dev/pytest>

<sup>3</sup><https://docs.python.org/3/library/pickle.html>

Dataset	#Samples	Repo context	Test case	Coverage (%)
CoNaLA Yin et al. [2018]	500	✗	✗	-
CONCODE Iyer et al. [2018]	2,000	✗	✗	-
HumanEval Chen et al. [2021]	164	✗	✓, <b>H</b>	99.43
MBPP Austin et al. [2021]	974	✗	✓, <b>H</b>	98.48
RepoCoder Zhang et al. [2023a]	373	✓	✓, <b>P</b>	-
CrossCodeEval Ding et al. [2023]	2,665	✓	✗	-
CoCoMIC Ding et al. [2024]	6,888	✓	✗	-
DevEval Li et al. [2024]	1,874	✓	✓, <b>P</b>	-
REPOEXEC	355	✓	✓, <b>A</b>	92.46
+ coverage-enhancement				96.25

Table 1: The comparison between popular code generation benchmarks and REPOEXEC. For test case, we denote **H**: Human annotated, **P**: Pre-existing, **A**: Automated.

### 3.3.2 Coverage Enhancement

Weak unit tests may allow incorrect implementations to pass Liu et al. [2023a]. To address this, we propose a strategy for enhancing test coverage using LLMs. Given the complexity of this task, requiring a strong understanding of code, we utilize GPT-3.5 to improve the quality of test cases. We provide GPT-3.5 with three prompts (see Appendix C) to handle challenging scenarios, including edge and corner cases. The initial generated tests serve as few-shot examples. We then extract the newly generated tests and ensure their correctness using our methodology from Section 3.3.1. Table 1 shows a coverage improvement of about 4% after enhancement, reaching 96.25%. The performance gap has also increased to over 5% (Appendix B), indicating greater robustness in the evaluation.

## 4 Data Characteristics

### 4.1 Data Formatting

Figure 2 illustrates the input data format used in REPOEXEC. We retain import information and append dependencies in the order presented in the import statements. The target function signature and natural language description are positioned at the end of the input prompt. We propose three prompt types with varying context lengths to test the reasoning capability of CodeLLMs in leveraging contexts for repo-level code generation:

- **Full-size context:** All contexts, including cross-file and in-file contexts, are preserved to assess the model’s ability to navigate and utilize the complete information available in the repository.
- **Medium-size context:** Class and function bodies are removed, while their signatures and docstrings are retained. This tests the model’s ability to infer the functionality and usage of dependencies based on their interfaces and documentation, reducing the input context length.
- **Small-size context:** Only the signatures of the dependencies are retained. This tests if CodeLLMs can infer the usage of dependencies in the target function given only the function signatures without docstrings, representing the most challenging case with minimal information.

Evaluating the model’s performance across these context sizes provides insights into the trade-offs between input context length and the model’s reasoning capabilities, helping to understand the optimal balance between providing sufficient information and minimizing input size for effective code generation at the repository level.

We follow Muennighoff et al. 2023 to define 2 types of prompt formats in our evaluation of REPOEXEC across LLMs: (1) BasePrompt, which concatenates all contexts with the target functions (Figure 2), and (2) InstructPrompt, which includes specific instructions for the LLMs to follow, utilizing two variations as input formats (further details and examples in Appendix D).

### 4.2 Dataset Stastistic

#### Comparison to Existing Benchmarks:

	#No problem		#No testcase		Avg tokens	
	Cross-file	Total	Avg Coverage (%)	Avg	Prompt	Solution
Full context	22.8%	355	96.25	98.38	362.92	78.46
Medium context	-	-	-	-	253.05	-
Short context	-	-	-	-	179.66	-

Table 2: Dataset attributes. Cross-file refers to the percentage of problems that involve cross-file dependencies. Tree-sitter is used for tokenization.

Table 1 compares the details of REPOEXEC with those of existing code generation benchmarks. Benchmarks that exclude execution-based evaluation Yin et al. [2018], Iyer et al. [2018], Ding et al. [2023, 2024] may gather substantial amounts of data; however, they are inadequate for assessing the quality of the generated code. For HumanEval and MBPP, the majority of problems involve standalone functions, which do not reflect real-world scenarios. Besides, benchmarks that rely on human-annotated and pre-existing test cases Chen et al. [2021], Austin et al. [2021], Zhang et al. [2023b], Li et al. [2024] are challenging to scale and control the test coverage. Finally, repository-context benchmarks (RepoCoder Zhang et al. [2023b], CrossCodeEval Ding et al. [2023], CoCoMIC Ding et al. [2024], and DevEval Li et al. [2024]) primarily focus on evaluating retrieval modules. In contrast, our approach emphasizes the generation module’s ability to understand and utilize human-provided dependency contexts.

**Dataset attributes:** Table 2 outlines the characteristics of REPOEXEC, including the total number of examples, the average number of test cases per problem, and the number of tokens in prompts and solutions.

## 5 Experiment

### 5.1 Evaluation metrics

We use the well-known *Pass@k* Chen et al. [2021] to measure the functional correctness of generation outputs. Furthermore, to assess the models’ ability to utilize provided dependencies in accordance with human intent, we introduce the *Dependency Invocation Rate (DIR)*. This metric represents the percentage of invoked dependencies out of the total number of dependencies provided. Let  $D_g$  denote the set of identifiers in the generated output, and  $D_s$  denote the set of provided dependencies extracted from the solution. The Dependency Invocation Rate (DIR) is calculated as follows:

$$\text{DIR} = \frac{|D_g \cap D_s|}{|D_s|}$$

*A higher DIR indicates that the model successfully incorporates a larger proportion of the provided dependencies into the generated code, demonstrating a better understanding of the dependencies’ relevance and their intended usage. Conversely, a lower DIR suggests that the model may struggle to identify and utilize the appropriate dependencies, potentially generating code that is less aligned with the human-written solution.*

### 5.2 Evaluation Results

We evaluated 13 CodeLLMs on REPOEXEC and presented the results (pass@1, pass@5, and DIR) in Table 3. We use nucleus sampling with temperature set to 0.2, top-p to 0.95, and 10 outputs generated for all models. The results show that retaining the full context of dependencies yields the best performance across all models. Surprisingly, Small-size context proves to be more effective than Medium-size context, which we attribute to the context’s input format using BasePrompt, potentially misleading the model into interpreting the dependency functions as few-shot examples. Using the Small-size context results in a fair decrease in performance compared to the Full context while effectively reducing the input length by half.

Codellama-34b-Python achieves the highest pass@1 rate among the models considered, and foundation models demonstrate greater effectiveness compared to instruction-tuned models on our dataset. However, our analysis reveals inherent limitations in both types of models:

Model		Full context			Medium context			Small context		
		pass@1	pass@5	DIR	pass@1	pass@5	DIR	pass@1	pass@5	DIR
BasePrompt										
Pre-models	CodeLlama-34b-Python Roziere et al. [2023]	42.93	49.54	68.85	35.92	42.95	58.15	39.80	45.79	64.23
	CodeLlama-13b-Python Roziere et al. [2023]	38.65	43.24	62.26	32.96	38.33	56.38	35.66	42.41	62.67
	StarCoder Li et al. [2023]	28.08	33.95	58.75	22.54	31.83	50.74	25.54	31.45	56.67
	StarCoder2-15b Lozhkov et al. [2024]	27.77	32.60	60.57	18.70	23.28	39.28	23.27	29.78	53.49
	Mixtral-8x7B-v0.1 Jiang et al. [2024]	22.82	29.14	55.90	19.38	25.25	47.22	20.54	26.30	53.40
	Phi-2 Javaheripi et al. [2023]	19.04	24.56	48.22	14.54	20.34	40.85	14.82	20.69	44.54
Inst-models	Phi-1 Gunasekar et al. [2023]	14.99	18.38	43.17	12.48	15.42	37.45	12.54	15.96	38.75
	WizardCoder-Python-13B-V1.0 Luo et al. [2023]	34.31	40.06	62.90	30.99	36.75	59.50	32.54	38.34	64.67
	Phind-CodeLlama-34B-v2	30.08	33.49	59.47	25.25	29.40	50.73	27.55	31.85	58.93
	CodeLlama-13b-Instruct Roziere et al. [2023]	28.56	32.67	57.09	26.25	30.72	49.48	26.73	33.50	54.53
	GPT-3.5	27.27	37.69	63.59	23.15	33.94	52.79	22.59	33.63	55.22
	DeepSeek-Coder-7b-Instruct Guo et al. [2024]	25.18	29.91	58.50	20.23	26.02	45.76	22.20	27.74	56.69
Mixtral-8x7B-Instruct-v0.1 Jiang et al. [2024]		23.41	28.71	59.83	19.04	24.83	52.75	20.45	25.84	58.17
InstructPrompt										
Inst-models	WizardCoder-Python-13B-V1.0	26.20	30.68	67.32	24.56	30.25	67.54	24.70	29.56	68.34
	CodeLlama-13b-Instruct	25.66	30.82	62.04	27.44	34.11	63.33	26.73	32.43	64.85
	GPT-3.5	23.82	39.10	40.55	19.62	36.03	37.48	19.00	34.00	35.28
	Mixtral-8x7B-Instruct-v0.1	18.11	23.04	67.73	18.54	23.12	69.66	15.38	20.61	68.86

Table 3: Pass@k (k= 1 and 5) and DIR results of various LLMs on REPOEXEC. **Bold scores** indicate the highest values, while Underlined scores represent the second highest. Pre- and Inst- denote Pretrained and Instruction-tuned, respectively.

Round	GPT-3.5	WizardCoder	CodeLlama-13b-Python
0	27.04	34.37	39.44
1	36.34	40.85	39.44
2	40.00	41.69	39.44
3	41.97	42.54	39.44

Table 4: Pass@1 scores of various models across three rounds of debugging. Round 0 represents the initial generation stage.

1. Instruction-tuned LLMs demonstrate a higher capacity for utilizing given dependencies than foundation LMs, sometimes enabling them to address edge or corner cases that foundation models have overlooked.
2. Despite the high DIR, instruction-tuning LMs may not utilize dependencies correctly and frequently produce overly complex code, leading to incorrect solutions.
3. Pretrained LLMs produce code that functions correctly but do not effectively reuse the provided dependencies, occasionally duplicating the implementation from the given context. This leads to redundancy and potentially creates technical debt or code smell issues [Maldonado and Shihab, 2015, Sierra et al., 2019, Rasool and Arshad, 2015, Santos et al., 2018].

These issues can lead to a low-quality codebase, requiring substantial human effort for reviewing and fixing, which may even exceed the effort needed to write the code from scratch. Details and examples are discussed further in Appendix E.

### 5.3 Generation with Multi-round Debugging

In this section, we examine the models’ self-refinement capabilities in enhancing generation performance. We provide the models with error output logs and ask them to fix the errors. We experiment with WizardCoder, GPT3.5, and CodeLlama-13b-Python. The number of rounds to debug is set to 3 and the input template is presented in Appendix F. In this experiment, we employ a greedy search algorithm to generate only a single output for each problem.

Table 4 shows the improvement across three rounds of debugging in various models. We observe that GPT-3.5 and WizardCoder demonstrate a high capacity for debugging with improvement of over 10% and 7% in pass@1, respectively, while CodeLlama fails to take advantage of this process. Additionally, the DIR has also shown a significant improvement (over 7%) after three rounds of debugging in these two instruction models (Figure 7). These findings indicate a promising approach



Model	Full context		Small context	
	Pass@1	DIR	Pass@1	DIR
StarCoder	28.08	58.67	25.49	56.67
StarCoder <sub>DepIT</sub>	<b>29.43</b>	<b>69.80</b>	<b>28.73</b>	<b>71.48</b>
StarCoder2	27.77	60.57	23.27	53.49
StarCoder2 <sub>DepIT</sub>	<b>28.45</b>	<b>69.76</b>	<b>27.27</b>	<b>73.98</b>
CodeLlama	38.65	62.26	35.66	62.67
CodeLlama <sub>DepIT</sub>	<b>38.85</b>	<b>68.89</b>	<b>36.93</b>	<b>73.19</b>

Table 5: Comparison of the performance of several models on REPOEXEC after instruction tuning for dependency calls (<sub>DepIT</sub>) with their pre-trained versions.

using self-refinement with debugging for code generation, which can enhance both the correctness and the utilization of given dependencies.

#### 5.4 Instruction-tuning with Code Dependencies

While the multi-round debugging experiment has demonstrated effectiveness in leveraging given dependencies to provide correct solutions (Section 5.3), it requires a strong model to generate good test cases and can be time-consuming due to the repeated generation and execution of code and test cases. To address these challenges, we propose an instruction-tuning dataset for fine-tuning base LLMs. We collected the 1,555 most-starred repositories from 2018 onward and extracted functions with their corresponding dependencies, following the procedure outlined in Section 3.2. We obtained 154,818 functions, of which 57,746 samples include docstrings. From this set, we utilized 50K samples with docstrings and applied instruction prompts, while 80K samples adhered to the raw code format (Full context). Recognizing the potential of the Small-context format, we allocated the remaining 20K samples to follow this structure. We fine-tuned StarCoder, StarCoder2, and CodeLlama-13b-Python models for 5 epochs with LoRa Hu et al. [2021] and used 10% of the training data as the validation set to select the best checkpoint.

Table 5 illustrates the efficacy of our training data. All 3 models demonstrate improvements in both Pass@1 and DIR after instruction tuning. Specifically, there is a slight increase of around 1% in Pass@k for all models, while DIR shows a significant improvement, reaching the highest scores (over 70%) compared to other models after tuning with our dataset. Notably, the performance improves significantly with the use of small context, achieving results comparable to those obtained with full context. This offers the potential for more efficient processing and reducing computational costs, allowing additional space to integrate various other types of context. In summary, our proposed instruction-tuning method can improve the model’s ability to both reuse dependencies and ensure the functional correctness of the output. Although self-refinement through multi-round debugging (Section 5.3) demonstrates greater effectiveness, instruction-tuning models only utilize single-turn generation, making them more efficient in practice. *To facilitate open research in this domain, we will publicly release this dataset.*

## 6 Conclusion

We introduce REPOEXEC, a novel Python code generation benchmark at the repository level with executable capabilities, designed to evaluate the alignment of generated code with developer intent and its correctness. Our experiments reveal that while pretrained language models (LLMs) excel in functional correctness, instruction-tuned models demonstrate proficiency in utilizing dependencies and debugging. However, existing models struggle to effectively reuse provided dependencies, potentially leading to technical debt and code smell issues. We provide a comprehensive evaluation of how different LLMs leverage code dependencies when generating code, and our findings show that the size of the context significantly impacts the final results.

We also introduce an instruction-tuning dataset that enhances dependency invocation accuracy and output correctness, even with limited context. This approach allows for the integration of additional

context types and mitigates large token length constraints. Our contributions establish a foundation for future research in code generation, providing valuable evaluation techniques and insights to drive the development of more capable and reliable models.

## 7 Limitations

REPOEXEC is constructed for evaluating models based on existing works. However, this approach can potentially lead to data leakage, especially when modern models are trained on similar datasets. If the benchmark relies heavily on known works, there’s a risk that the model may inadvertently learn from specific patterns or features present in those works, compromising its generalization ability. Despite this concern, experimental results demonstrate that even modern models struggle to handle the challenges posed by REPOEXEC, indicating that the benchmark remains a valuable tool for assessing model performance. However, employing new repositories can help mitigate these phenomena.

In this work, we only consider one level of dependency context, specifically the dependencies directly called from the target function. While this simplification facilitates manageable analysis and model development, it may not fully capture the valuable context necessary for leveraging models effectively. However, incorporating deep dependencies could significantly extend the input length, posing challenges in managing long context inputs and potentially exceeding the maximum input length. Our approach has demonstrated promising outcomes with the short context version, creating opportunities for integrating additional input context. Future research could explore incorporating multiple levels of dependencies, creating a more comprehensive graph that includes transitive dependencies, indirect calls, and broader contextual information. By doing so, we could enhance the model’s understanding of code interactions and improve its ability to handle intricate software execution scenarios.

## References

- Loubna Ben Allal, Raymond Li, Denis Kocetkov, Chenghao Mou, Christopher Akiki, Carlos Munoz Ferrandis, Niklas Muennighoff, Mayank Mishra, Alex Gu, Manan Dey, et al. Santacoder: don’t reach for the stars! *arXiv preprint arXiv:2301.03988*, 2023.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Nghi DQ Bui, Hung Le, Yue Wang, Junnan Li, Akhilesh Deepak Gotmare, and Steven CH Hoi. Codetf: One-stop transformer library for state-of-the-art code llm. *arXiv preprint arXiv:2306.00029*, 2023.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Yanguibo Ding, Zijian Wang, Wasi Uddin Ahmad, Hantian Ding, Ming Tan, Nihal Jain, Murali Krishna Ramanathan, Ramesh Nallapati, Parminder Bhatia, Dan Roth, and Bing Xiang. Crosscodeeval: A diverse and multilingual benchmark for cross-file code completion. In Alice Oh, Tristan Naumann, Amir Globerson, Kate Saenko, Moritz Hardt, and Sergey Levine, editors, *Advances in Neural Information Processing Systems 36: Annual Conference on Neural Information Processing Systems 2023, NeurIPS 2023, New Orleans, LA, USA, December 10 - 16, 2023*, 2023.
- Yanguibo Ding, Zijian Wang, Wasi U. Ahmad, Murali Krishna Ramanathan, Ramesh Nallapati, Parminder Bhatia, Dan Roth, and Bing Xiang. CoCoMIC: Code completion by jointly modeling in-file and cross-file context. In Nicoletta Calzolari, Min-Yen Kan, Veronique Hoste, Alessandro Lenci, Sakriani Sakti, and Nianwen Xue, editors, *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, May 2024.
- Zhangyin Feng, Daya Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Daxin Jiang, et al. Codebert: A pre-trained model for programming and natural languages. *arXiv preprint arXiv:2002.08155*, 2020.

- Suriya Gunasekar, Yi Zhang, Jyoti Aneja, Caio César Teodoro Mendes, Allie Del Giorno, Sivakanth Gopi, Mojan Javaheripi, Piero Kauffmann, Gustavo de Rosa, Olli Saarikivi, et al. Textbooks are all you need. *arXiv preprint arXiv:2306.11644*, 2023.
- Daya Guo, Qihao Zhu, Dejian Yang, Zhenda Xie, Kai Dong, Wentao Zhang, Guanting Chen, Xiao Bi, Y Wu, YK Li, et al. Deepseek-coder: When the large language model meets programming—the rise of code intelligence. *arXiv preprint arXiv:2401.14196*, 2024.
- Rajarshi Haldar and Julia Hockenmaier. Analyzing the performance of large language models on code summarization. *arXiv preprint arXiv:2404.08018*, 2024.
- Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with apps. *NeurIPS*, 2021.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- Srinivasan Iyer, Ioannis Konstas, Alvin Cheung, and Luke Zettlemoyer. Mapping language to code in programmatic context. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 1643–1652, 2018.
- Mojan Javaheripi, Sébastien Bubeck, Marah Abdin, Jyoti Aneja, Sebastien Bubeck, Caio César Teodoro Mendes, Weizhu Chen, Allie Del Giorno, Ronen Eldan, Sivakanth Gopi, et al. Phi-2: The surprising power of small language models. *Microsoft Research Blog*, 2023.
- Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024.
- Caroline Lemieux, Jeevana Priya Inala, Shuvendu K Lahiri, and Siddhartha Sen. Codamosa: Escaping coverage plateaus in test generation with pre-trained large language models. In *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*, pages 919–931. IEEE, 2023.
- Jia Li, Ge Li, Yunfei Zhao, Yongmin Li, Huanyu Liu, Hao Zhu, Lecheng Wang, Kaibo Liu, Zheng Fang, Lanshen Wang, et al. Deval: A manually-annotated code generation benchmark aligned with real-world code repositories. *arXiv preprint arXiv:2405.19856*, 2024.
- Raymond Li, Loubna Ben Allal, Yangtian Zi, Niklas Muennighoff, Denis Kocetkov, Chenghao Mou, Marc Marone, Christopher Akiki, Jia Li, Jenny Chim, et al. Starcoder: may the source be with you! *arXiv preprint arXiv:2305.06161*, 2023.
- Dianshu Liao, Shidong Pan, Qing Huang, Xiaoxue Ren, Zhenchang Xing, Huan Jin, and Qinying Li. Context-aware code generation framework for code repositories: Local, global, and third-party library awareness. *arXiv preprint arXiv:2312.05772*, 2023.
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and LINGMING ZHANG. Is your code generated by chatGPT really correct? rigorous evaluation of large language models for code generation. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023a. URL <https://openreview.net/forum?id=1qv610Cu7>.
- Tianyang Liu, Canwen Xu, and Julian McAuley. Repobench: Benchmarking repository-level code auto-completion systems. *arXiv preprint arXiv:2306.03091*, 2023b.
- Anton Lozhkov, Raymond Li, Loubna Ben Allal, Federico Cassano, Joel Lamy-Poirier, Nouamane Tazi, Ao Tang, Dmytro Pykhtar, Jiawei Liu, Yuxiang Wei, et al. Starcoder 2 and the stack v2: The next generation. *arXiv preprint arXiv:2402.19173*, 2024.
- Ziyang Luo, Can Xu, Pu Zhao, Qingfeng Sun, Xiubo Geng, Wenxiang Hu, Chongyang Tao, Jing Ma, Qingwei Lin, and Daxin Jiang. Wizardcoder: Empowering code large language models with evol-instruct. *arXiv preprint arXiv:2306.08568*, 2023.

- Everton da S Maldonado and Emad Shihab. Detecting and quantifying different types of self-admitted technical debt. In *2015 IEEE 7Th international workshop on managing technical debt (MTD)*, pages 9–15. IEEE, 2015.
- Dung Nguyen Manh, Nam Le Hai, Anh T. V. Dau, Anh Minh Nguyen, Khanh Nghiem, Jin Guo, and Nghi D. Q. Bui. The vault: A comprehensive multilingual dataset for advancing code understanding and generation. In Liling Tan, Dmitrijs Milajevs, Geeticka Chauhan, Jeremy Gwinnup, and Elijah Rippeth, editors, *Proceedings of the 3rd Workshop for Natural Language Processing Open Source Software (NLP-OSS 2023)*, pages 219–244, Singapore, December 2023. Association for Computational Linguistics.
- Debanjan Mondal, Abhilasha Lodha, Ankita Sahoo, and Beena Kumari. Robust code summarization. In Dieuwke Hupkes, Verna Dankers, Khuyagbaatar Batsuren, Koustuv Sinha, Amirhossein Kazemnejad, Christos Christodoulopoulos, Ryan Cotterell, and Elia Bruni, editors, *Proceedings of the 1st GenBench Workshop on (Benchmarking) Generalisation in NLP*, pages 65–75, Singapore, December 2023. Association for Computational Linguistics.
- Niklas Muennighoff, Qian Liu, Armel Zebaze, Qinkai Zheng, Binyuan Hui, Terry Yue Zhuo, Swayam Singh, Xiangru Tang, Leandro Von Werra, and Shayne Longpre. Octopack: Instruction tuning code large language models. *arXiv preprint arXiv:2308.07124*, 2023.
- Erik Nijkamp, Bo Pang, Hiroaki Hayashi, Lifu Tu, Huan Wang, Yingbo Zhou, Silvio Savarese, and Caiming Xiong. Codegen: An open large language model for code with multi-turn program synthesis. *arXiv preprint arXiv:2203.13474*, 2022.
- Nikhil Pinnaparaju, Reshinth Adithyan, Duy Phung, Jonathan Tow, James Baicoianu, Ashish Datta, Maksym Zhuravinskiy, Dakota Mahan, Marco Bellagente, Carlos Riquelme, et al. Stable code technical report. *arXiv preprint arXiv:2404.01226*, 2024.
- Ghulam Rasool and Zeeshan Arshad. A review of code smell mining techniques. *Journal of Software: Evolution and Process*, 27(11):867–895, 2015.
- Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*, 2023.
- José Amancio M Santos, João B Rocha-Junior, Luciana Carla Lins Prates, Rogeres Santos Do Nascimento, Mydiã Falcão Freitas, and Manoel Gomes De Mendonça. A systematic review on the code smell effect. *Journal of Systems and Software*, 144:450–477, 2018.
- Disha Shrivastava, Hugo Larochelle, and Daniel Tarlow. Repository-level prompt generation for large language models of code. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *International Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, volume 202 of *Proceedings of Machine Learning Research*, pages 31693–31715. PMLR, 2023.
- Giancarlo Sierra, Emad Shihab, and Yasutaka Kamei. A survey of self-admitted technical debt. *Journal of Systems and Software*, 152:70–82, 2019.
- Ankita Nandkishor Sontakke, Manasi Patwardhan, Lovekesh Vig, Raveendra Kumar Medicherla, Ravindra Naik, and Gautam Shroff. Code summarization: Do transformers really understand code? In *Deep Learning for Code Workshop*, 2022.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shrutu Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.
- Yue Wang, Weishi Wang, Shafiq Joty, and Steven CH Hoi. Codet5: Identifier-aware unified pre-trained encoder-decoder models for code understanding and generation. *arXiv preprint arXiv:2109.00859*, 2021.

- Yue Wang, Hung Le, Akhilesh Deepak Gotmare, Nghi DQ Bui, Junnan Li, and Steven CH Hoi. Codet5+: Open code large language models for code understanding and generation. *arXiv preprint arXiv:2305.07922*, 2023.
- Frank F Xu, Uri Alon, Graham Neubig, and Vincent Josua Hellendoorn. A systematic evaluation of large language models of code. In *Proceedings of the 6th ACM SIGPLAN International Symposium on Machine Programming*, pages 1–10, 2022.
- Pengcheng Yin, Bowen Deng, Edgar Chen, Bogdan Vasilescu, and Graham Neubig. Learning to mine aligned code and natural language pairs from stack overflow. In *Proceedings of the 15th international conference on mining software repositories*, pages 476–486, 2018.
- Fengji Zhang, Bei Chen, Yue Zhang, Jacky Keung, Jin Liu, Daoguang Zan, Yi Mao, Jian-Guang Lou, and Weizhu Chen. Repocoder: Repository-level code completion through iterative retrieval and generation. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP 2023, Singapore, December 6-10, 2023*, pages 2471–2484. Association for Computational Linguistics, 2023a.
- Fengji Zhang, Bei Chen, Yue Zhang, Jacky Keung, Jin Liu, Daoguang Zan, Yi Mao, Jian-Guang Lou, and Weizhu Chen. Repocoder: Repository-level code completion through iterative retrieval and generation. *arXiv preprint arXiv:2303.12570*, 2023b.
- Kechi Zhang, Jia Li, Ge Li, Xianjie Shi, and Zhi Jin. Codeagent: Enhancing code generation with tool-integrated agent systems for real-world repo-level coding challenges. *arXiv preprint arXiv:2401.07339*, 2024.
- Tianyu Zheng, Ge Zhang, Tianhao Shen, Xueling Liu, Bill Yuchen Lin, Jie Fu, Wenhui Chen, and Xiang Yue. Opencodeinterpreter: Integrating code generation with execution and refinement. *arXiv preprint arXiv:2402.14658*, 2024.

# Appendix

## A Evaluation Metrics: Match-Based vs. Execution-Based

Several code generation benchmarks utilize match-based metrics like Edit similarity (ES), BLEU, and CodeBLEU for evaluation [Ding et al., 2024, Shrivastava et al., 2023, Ding et al., 2023, Liao et al., 2023, Yin et al., 2018, Iyer et al., 2018]. These metrics are straightforward to apply and may exhibit a strong correlation with execution metrics such as Pass@k. However, they cannot accurately measure functional correctness. For instance, comparing two Python code snippets where the only difference is the ":" character could result in a high ES and BLEU score. Nevertheless, one snippet may contain a syntax error, highlighting a limitation in these metrics for assessing true functionality.




Figure 3: Correlation between Match-based metrics and Execution-based metric (Pass@1).

Figure 3 demonstrates that these metrics on average can achieve a strong correlation with Pass@k, with CodeBLEU showing the highest correlation with Pearson score of 0.92. However, upon closer inspection of the score distribution between correct and incorrect solutions (Figure 4), a considerable overlap becomes apparent. This underscores the limitation of match-based metrics in accurately measuring the correctness of code generation.




Figure 4: Match-based metric distributions between Correct and Incorrect solutions

## B Coverage Enhancement Effectiveness

Weak unit tests may inadvertently allow incorrect implementations to be determined as correct. Even with human-written tests, the overlooked coverage rates lead to evaluations that are incomplete and potentially misleading Liu et al. [2023a]. We present evidence supporting this argument, underscoring the limitations of prior work on code generation within repository-level contexts. As depicted in Figure 5, enhancing the number of test cases and coverage rates leads to a significant increase in the identification of incorrect generated solutions, causing the Pass@1 score to drop markedly (by

over 5%). We investigated several solutions and found that most of the generated results did not fully utilize the given context (considered as human-provided). Instead, they primarily focused on addressing the problem described in the given natural language description. This indirectly overlooks the developer’s intentions, such as testing edge or corner cases, highlighting the limitations in following and understanding the provided intent and dependency context in these models. In summary, these findings underscore the effectiveness and importance of maintaining high-quality test cases for evaluation purposes.




Figure 5: Performance of various CodeLMs on REPOEXEC before (bf-) and after (af-) Coverage Enhancement (CovEn) stage.

## C Test case generation

### Initial Test Generation Prompt

```
{function_under_test}
# test to check the correctness of "{function_name}" function
assert
```

## Coverage Enhancement Prompts

### **Prompt 1:**

Here are some Python unit test functions and the focal function that they test:

# Test functions:

{existing\_test\_functions}

# Focal function:

{function\_under\_test}

Write more unit test functions that will increase the test coverage of the function under test.

---

### **Prompt 2:**

Here are some Python unit test functions and the focal function that they test:

# Test functions:

{existing\_test\_functions}

# Focal function:

{function\_under\_test}

Write more unit test functions that will cover corner cases missed by the original and will increase the test coverage of the function under test.

---

### **Prompt 3:**

Here is a focal function under test:

{function\_under\_test}

This function under test can be tested with these Python unit test functions:

{existing\_test\_functions}

Here is an extended version of the unit test function that includes additional unit test cases that will cover methods, edge cases, corner cases, and other features of the function under test that were missed by the original unit test functions:





## D Data Formating

### D.1 BasePrompt

#### Example of Full Context

```
import base64
import random
import unicodedata
import zlib
from typing import Union
from uuid import uuid4
from ._regex import *
from .errors import InvalidInputError
from .validation import is_snake_case, is_full_string, is_camel_case, is_integer,
    is_string

CAMEL_CASE_REPLACE_RE = re.compile(r'([a-z]|[A-Z]+)(?=[A-Z])')

class InvalidInputError(TypeError):
    """
    Custom error raised when received object is not a string as expected.
    """

    def __init__(self, input_data: Any):
        """
        :param input_data: Any received object
        """
        type_name = type(input_data).__name__
        msg = 'Expected "str", received "{}'.format(type_name)
        super().__init__(msg)

def is_string(obj: Any) -> bool:
    """
    Checks if an object is a string.

    *Example:*

    >>> is_string('foo') # returns true
    >>> is_string(b'foo') # returns false

    :param obj: Object to test.
    :return: True if string, false otherwise.
    """
    return isinstance(obj, str)

def is_camel_case(input_string: Any) -> bool:
    """
    Checks if a string is formatted as camel case.

    A string is considered camel case when:

    - it's composed only by letters ([a-zA-Z]) and optionally numbers ([0-9])
    - it contains both lowercase and uppercase letters
    - it does not start with a number

    *Examples:*

    >>> is_camel_case('MyString') # returns true
    >>> is_camel_case('mystring') # returns false

    :param input_string: String to test.
    :type input_string: str
    :return: True for a camel case string, false otherwise.
    """
    return is_full_string(input_string) and CAMEL_CASE_TEST_RE.match(input_string) is
    not None

def camel_case_to_snake(input_string, separator='_'):
    """
    Convert a camel case string into a snake case one.
    (The original string is returned if is not a valid camel case string)

    *Example:*

    >>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
    this_is_a_camel_case_string_test'

    :param input_string: String to convert.
    :type input_string: str
    :param separator: Sign to use as separator.
    :type separator: str
    :return: Converted string.
    """
```

## Example of Medium Context

```
import base64
import random
import unicodedata
import zlib
from typing import Union
from uuid import uuid4
from .regex import *
from .errors import InvalidInputError
from .validation import is_snake_case, is_full_string, is_camel_case, is_integer,
    is_string

CAMEL_CASE_REPLACE_RE = re.compile(r'([a-z]|[A-Z]+)(?=[A-Z])')

class InvalidInputError(TypeError):
    """
    Custom error raised when received object is not a string as expected.
    """

    def __init__(self, input_data: Any):
        """
        :param input_data: Any received object
        """

def is_string(obj: Any) -> bool:
    """
    Checks if an object is a string.

    *Example:*

    >>> is_string('foo') # returns true
    >>> is_string(b'foo') # returns false

    :param obj: Object to test.
    :return: True if string, false otherwise.
    """

def is_camel_case(input_string: Any) -> bool:
    """
    Checks if a string is formatted as camel case.

    A string is considered camel case when:

    - it's composed only by letters ([a-zA-Z]) and optionally numbers ([0-9])
    - it contains both lowercase and uppercase letters
    - it does not start with a number

    *Examples:*

    >>> is_camel_case('MyString') # returns true
    >>> is_camel_case('mystring') # returns false

    :param input_string: String to test.
    :type input_string: str
    :return: True for a camel case string, false otherwise.
    """

def camel_case_to_snake(input_string, separator='_'):
    """
    Convert a camel case string into a snake case one.
    (The original string is returned if is not a valid camel case string)

    *Example:*

    >>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
    this_is_a_camel_case_string_test'

    :param input_string: String to convert.
    :type input_string: str
    :param separator: Sign to use as separator.
    :type separator: str
    :return: Converted string.
    """
```

## Example of Short Context

```
import base64
import random
import unicodedata
import zlib
from typing import Union
from uuid import uuid4
from ._regex import *
from .errors import InvalidInputError
from .validation import is_snake_case, is_full_string, is_camel_case, is_integer,
    is_string

CAMEL_CASE_REPLACE_RE = re.compile(r'([a-z]|[A-Z]+)(?=[A-Z])')

class InvalidInputError(TypeError):
    def __init__(self, input_data: Any):

def is_string(obj: Any) -> bool:

def is_camel_case(input_string: Any) -> bool:

def camel_case_to_snake(input_string, separator='_'):
    """
    Convert a camel case string into a snake case one.
    (The original string is returned if is not a valid camel case string)

    *Example:*

    >>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
    this_is_a_camel_case_string_test'

    :param input_string: String to convert.
    :type input_string: str
    :param separator: Sign to use as separator.
    :type separator: str
    :return: Converted string.
    """
```

## D.2 InstructPrompt

### Instruction Prompt Templates

#### **Prompt 1:**

### Instruction:

Write a Python function `{target\_function\_signature}` to solve the following problem:  
{target\_function\_docstring}

### Response:

{BasePrompt}

---

#### **Prompt 2:**

### Instruction:

{dependency\_context}  
The provided code snippet includes necessary dependencies for implementing the `{target\_function\_name}` function. Write a Python function `{target\_function\_signature}` to solve the following problem:  
{target\_function\_docstring}

### Response:

{target\_function\_prompt}

## Example of Prompt 1 for Short Context

```
### Instruction:

Write a Python function `camel_case_to_snake(input_string, separator='_')` to solve the
following problem:
"""
Convert a camel case string into a snake case one.
(The original string is returned if is not a valid camel case string)

*Example:*

>>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
this_is_a_camel_case_string_test'

:param input_string: String to convert.
:type input_string: str
:param separator: Sign to use as separator.
:type separator: str
:return: Converted string.
"""

### Response:
import base64
import random
import unicodedata
import zlib
from typing import Union
from uuid import uuid4
from ._regex import *
from .errors import InvalidInputError
from .validation import is_snake_case, is_full_string, is_camel_case, is_integer,
is_string

CAMEL_CASE_REPLACE_RE = re.compile(r'([a-z]|[A-Z])+(?=[A-Z])')

class InvalidInputError(TypeError):
    def __init__(self, input_data: Any):

def is_string(obj: Any) -> bool:

def is_camel_case(input_string: Any) -> bool:

def camel_case_to_snake(input_string, separator='_'):
    """
    Convert a camel case string into a snake case one.
    (The original string is returned if is not a valid camel case string)

    *Example:*

    >>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
    this_is_a_camel_case_string_test'

    :param input_string: String to convert.
    :type input_string: str
    :param separator: Sign to use as separator.
    :type separator: str
    :return: Converted string.
    """
```

### Example of Prompt 2 for Short Context

```
### Instruction

import base64
import random
import unicodedata
import zlib
from typing import Union
from uuid import uuid4
from ._regex import *
from .errors import InvalidInputError
from .validation import is_snake_case, is_full_string, is_camel_case, is_integer,
    is_string

CAMEL_CASE_REPLACE_RE = re.compile(r'([a-z]|[A-Z]+)(?=[A-Z])')

class InvalidInputError(TypeError):
    def __init__(self, input_data: Any):

def is_string(obj: Any) -> bool:

def is_camel_case(input_string: Any) -> bool:

The provided code snippet includes necessary dependencies for implementing the `
camel_case_to_snake` function. Write a Python function `camel_case_to_snake(
    input_string, separator='_')` to solve the following problem:
"""
Convert a camel case string into a snake case one.
(The original string is returned if is not a valid camel case string)

*Example:*

>>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
    this_is_a_camel_case_string_test'

:param input_string: String to convert.
:type input_string: str
:param separator: Sign to use as separator.
:type separator: str
:return: Converted string.
"""

### Response:
def camel_case_to_snake(input_string, separator='_'):
    """
    Convert a camel case string into a snake case one.
    (The original string is returned if is not a valid camel case string)

    *Example:*

    >>> camel_case_to_snake('ThisIsACamelStringTest') # returns '
        this_is_a_camel_case_string_test'

    :param input_string: String to convert.
    :type input_string: str
    :param separator: Sign to use as separator.
    :type separator: str
    :return: Converted string.
    """
```

## E Studied LLMs: Supplemental results

In this section, we offer supplementary results from the evaluation of LLMs on REPOEXEC. Table 3 presents the Dependency Invocation Rate (DIR) for the experimented LLMs. When comparing models of the same size, it is shown that instruction-tuned models more effectively follow human intent in utilizing the provided dependencies with InstructPrompt. For example, WizardCoder outperforms CodeLlama by 5%, and the instruction-tuned version of Mixtral-8x7B shows a 10% improvement over its foundation version. This highlights the strong capability of instruction-tuned models to follow the given context effectively. Besides, using the Medium context leads to a significant decline in both




Figure 6: Percentage of generated outputs that result in empty functions across various context types.

Pass@k and DIR using BasePrompt. This implies the generation of empty function bodies using this template.

Indeed, Figure 6 illustrates the proportion of generated functions that are empty for each LLM using BasePrompt. The findings indicate that utilizing Medium context results in a substantial number of empty functions. This may be due to the input format of the context when using BasePrompt, which can mislead the model into interpreting dependency functions as few-shot examples. In the Medium context, the function bodies of dependencies are removed, making their format identical to the target function prompt. This similarity can mislead the LMs, resulting in empty solutions. Particularly, Starcoder-2 is heavily impacted by this issue, as over 31% of its generated results are empty functions, revealing a significant weakness of the model. Meanwhile, short context effectively decreases the occurrence of empty function generation by the model and, in certain instances, improves models’ ability in dependency calls (e.g. CodeLlama-13b-Python, WizardCoder-Python-13B-V1.0, and Mixtral-8x7B-Instruct-v0.1 in Table 3). We believe that the following reasons could contribute to this observation. Firstly, employing short context reduces the input token count, preventing truncation when exceeding the maximum length limit, thus allowing uninterrupted solution generation by the model. This reduced context enables models to concentrate exclusively on dependency signatures, thereby enhancing the probability that generated solutions effectively utilize these dependency token names. Moreover, function names hold substantial semantic value by delineating the function’s purpose. Many studies have underscored that code summarization heavily relies on extracting information from function names [Haldar and Hockenmaier, 2024, Mondal et al., 2023, Sontakke et al., 2022]. Therefore, this concise representation of dependencies has the potential to improve how models utilize dependencies in generating code.

Additionally, Table 6 presents examples that support our findings in Section 5.2. For instance, in the first example, we observe that the instruction-tuned model can effectively utilize the given dependencies to manage edge cases, whereas the pretrained model fails to do so. This supports our initial findings. Meanwhile, the second and third examples demonstrate that pretrained models often reimplement or devise workarounds instead of leveraging the available context. Besides, in the third example, the instruction-tuned model correctly identifies the relevant case but fails to generate the correct solution. This may suggest that hallucinations complicate the outputs generated by instruction-tuned models.



## F Multi-round Debugging

We employ Multi-round debugging in code generation, which iteratively refines and improves the generated code through multiple cycles of debugging. Following the execution of unit tests on the generated functions, we extract the error log if the code fails to run. We employ the following prompt to utilize the model for bug fixing. This process is iterated multiple times until either the correct code is achieved or the maximum number of rounds is reached. Specifically, we set the maximum number of rounds to 3 and experimented on three models WizardCoder, GPT-3.5 and CodeLlama-13b-Python.

### Prompt for Debugging

```
{dependency_context}
# The provided code snippet includes necessary dependencies for
  implementing the `{target_function_name}` function. Write a
  Python function `{target_function_signature}` to solve the
  following problem:
{target_function_docstring}

# Here is the current solution.
{error_solution}

# When executing the below test case.
{failed_test_case}

# The provided python code solution fails the test with the
  following errors, please correct them.
{error_log}

# Please provide the modified code for me to review and provide
  feedback.
{target_function_prompt}
```

Table 4 shows the improvement across three rounds of debugging in various models. We observe that GPT-3.5 and WizardCoder demonstrate a high capacity for debugging with improvement of over 10% and 7% in Pass@1, respectively, while CodeLlama fails to take advantage of this process. Additionally, the DIR has also shown a significant improvement (over 7%) after three rounds of debugging in these two instruction models (Figure 7). These findings indicate a promising approach using self-refinement with debugging for code generation, which can enhance both the correctness and the utilization of given dependencies.

We also present data on the number of error types corrected in each round of WizardCoder, as illustrated in Figure 8. We can see that AssertionError makes up the majority of errors across all rounds. This error type indicates either incorrect outputs from the generated code or the presence of empty function bodies (return None). However, by incorporating the test output guide, the model effectively addressed most of these errors. Furthermore, fundamental issues like SyntaxError or AttributeError were promptly rectified during the initial round.




Figure 7: Improvement of instruction-tuning models on Pass@1 and DIR after 3-round debugging process.




Figure 8: Fixed error types of WizardCoder across 3 rounds of the debugging process.

Input Prompt & Solution	Generation output
-------------------------	-------------------

<pre> # Prompt: import base64 import random import unicodedata ...  class InvalidInputError(TypeError):     ...  def is_string(obj: Any) -&gt; bool:     ...  def reverse(input_string: str) -&gt; str:     """     Returns the string with its chars reversed.      *Example:*      &gt;&gt;&gt; reverse('hello') # returns 'olleh'      :param input_string: String to revert.     :type input_string: str     :return: Reversed string.     """ </pre>	<pre> # Model type: Pretrained # Passed: No def reverse(input_string: str) -&gt; str:     """     Returns the string with its chars reversed.      *Example:*      &gt;&gt;&gt; reverse('hello') # returns 'olleh'      :param input_string: String to revert.     :type input_string: str     :return: Reversed string.     """     return input_string[::-1] </pre>
<pre> # Solution: def reverse(input_string: str) -&gt; str:     """     Returns the string with its chars reversed.      *Example:*      &gt;&gt;&gt; reverse('hello') # returns 'olleh'      :param input_string: String to revert.     :type input_string: str     :return: Reversed string.     """     if not is_string(input_string):         raise InvalidInputError(input_string)      return input_string[::-1] </pre>	<pre> # Model type: Instruction-tuned # Passed: Yes def reverse(input_string: str) -&gt; str:     """     Returns the string with its chars reversed.      *Example:*      &gt;&gt;&gt; reverse('hello') # returns 'olleh'      :param input_string: String to revert.     :type input_string: str     :return: Reversed string.     """     if not is_string(input_string):         raise InvalidInputError(input_string)     return input_string[::-1] </pre>

<pre> # Prompt: import random import luigi import luigi.configuration ...  class hadoopcli(luigi.Config):     command = ...     version = ...  def get_configured_hadoop_version():     """     CDH4 (hadoop 2+) has a slightly different     syntax for interacting with hdfs     via the command line.      The default version is CDH4, but one can     override     this setting with "cdh3" or "apache1" in     the hadoop section of the config     in order to use the old syntax.     """ </pre>	<pre> # Model type: Pretrained # Passed: Yes def get_configured_hdfs_client():     """     This is a helper that fetches the     configuration value for 'client' in     the [hdfs] section. It will return the     client that retains backwards     compatibility when 'client' isn't     configured.     """     return luigi.configuration.get_config().get ('hdfs', 'client', 'hadoopcli') </pre>
<pre> # Solution: def get_configured_hadoop_version():     """     CDH4 (hadoop 2+) has a slightly different     syntax for interacting with hdfs     via the command line.      The default version is CDH4, but one can     override     this setting with "cdh3" or "apache1" in     the hadoop section of the config     in order to use the old syntax.     """     return hadoopcli().version.lower() </pre>	<pre> # Model type: Instruction-tuned # Passed: No def get_configured_hadoop_version():     """     CDH4 (hadoop 2+) has a slightly different     syntax for interacting with hdfs     via the command line.      The default version is CDH4, but one can     override     this setting with "cdh3" or "apache1" in     the hadoop section of the config     in order to use the old syntax.     """     hadoop_cli = hadoopcli().command.split()     hadoop_version = hadoopcli().version     if hadoop_version == "cdh3":         return hadoop_cli + ["fs"]     elif hadoop_version == "apache1":         return hadoop_cli + ["dfs"]     else:         return hadoop_cli </pre>

<pre> # Prompt: import asyncio import builtins import collections ...  def is_future(x: Any) -&gt; bool:     return isinstance(x, FUTURES)  def _create_future() -&gt; Future:     future = Future() # type: Future     ...  def maybe_future(x: Any) -&gt; Future:     """Converts ``x`` into a ``Future``.      If ``x`` is already a ``Future``, it is     simply returned; otherwise     it is wrapped in a new ``Future``. This is     suitable for use as     ``result = yield gen.maybe_future(f())``     when you don't know whether     ``f()`` returns a ``Future`` or not.      .. deprecated:: 4.3     This function only handles ``Futures``, not     other yieldable objects.     Instead of ``maybe_future``, check for the     non-future result types     you expect (often just ``None``), and ``     yield`` anything unknown.     """ </pre>	<pre> # Model type: Pretrained # Passed: Yes def maybe_future(x: Any) -&gt; Future:     """Converts ``x`` into a ``Future``.      If ``x`` is already a ``Future``, it is     simply returned; otherwise     it is wrapped in a new ``Future``. This is     suitable for use as     ``result = yield gen.maybe_future(f())``     when you don't know whether     ``f()`` returns a ``Future`` or not.      .. deprecated:: 4.3     This function only handles ``Futures``, not     other yieldable objects.     Instead of ``maybe_future``, check for the     non-future result types     you expect (often just ``None``), and ``     yield`` anything unknown.     """      if is_future(x):         return x     future = Future() # type: Future     future.set_result(x)     return future </pre>
<pre> # Solution: def maybe_future(x: Any) -&gt; Future:     """Converts ``x`` into a ``Future``.      If ``x`` is already a ``Future``, it is     simply returned; otherwise     it is wrapped in a new ``Future``. This is     suitable for use as     ``result = yield gen.maybe_future(f())``     when you don't know whether     ``f()`` returns a ``Future`` or not.      .. deprecated:: 4.3     This function only handles ``Futures``,     not other yieldable objects.     Instead of ``maybe_future``, check for the     non-future result types     you expect (often just ``None``), and ``     yield`` anything unknown.     """      if is_future(x):         return x     else:         fut = _create_future()         fut.set_result(x)         return fut </pre>	<pre> # Model type: Instruction-tuned # Passed: Yes def maybe_future(x: Any) -&gt; Future:     """Converts ``x`` into a ``Future``.      If ``x`` is already a ``Future``, it is     simply returned; otherwise     it is wrapped in a new ``Future``. This is     suitable for use as     ``result = yield gen.maybe_future(f())``     when you don't know whether     ``f()`` returns a ``Future`` or not.      .. deprecated:: 4.3     This function only handles ``Futures``, not     other yieldable objects.     Instead of ``maybe_future``, check for the     non-future result types     you expect (often just ``None``), and ``     yield`` anything unknown.     """      if isinstance(x, Future):         return x     elif isinstance(x, Awaitable):         return asyncio.ensure_future(x)     else:         future = Future()         future.set_result(x)         return future </pre>

Table 6: Output samples of Pretrained and Instruction-tuned models.